

Risk Management process

1. Risk analysis
 - Analyze the frequency and severity of harms
2. Risk evaluation
 - Determine what is acceptable risk
3. Risk control
 - How to prevent hazard and/or reduce harm?
4. Residual risk acceptability
 - Are there new risks introduced by control measures?

Ethiopian airline crash

- Flight ET302 from Addis Ababa for Nairobi
- Crashed on March 10, 2019
- Crashed 6 mins after takeoff
- All 157 people on board died
- Boeing 737-Max 8




Similar Crash

- Lion Airline JT610 Oct 29, 2018
- Crashed 12 mins after takeoff
- All 189 people on board died






Boeing 737 Max 8/9

- The Max's engines were bigger and mounted farther upward on its wings




Stall

- Angle of attack (AOA): the angle between velocity vector and a line along the fuselage
- High AOA results in stall, which leads to crash
- The new engine configuration of 737 max makes it more likely to have high AOA

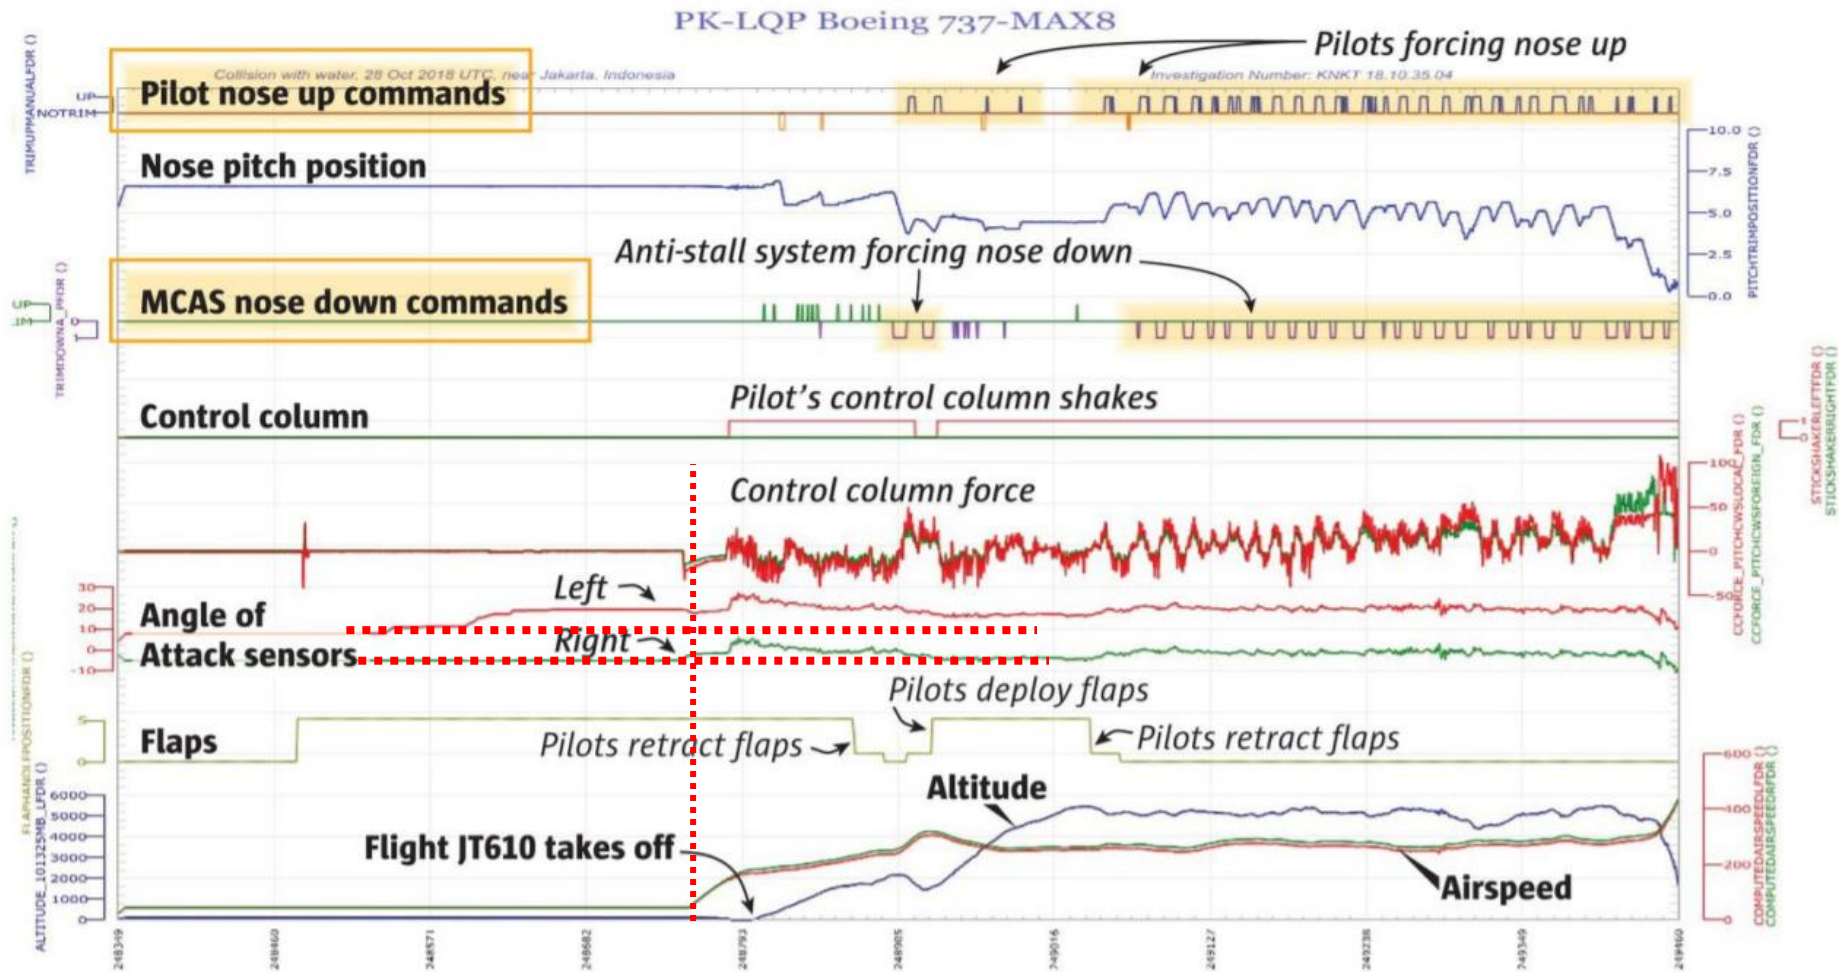


Maneuvering Characteristics Augmentation System (MCAS)

- MCAS to automatically push the nose down after detecting high AOA
- Two AOA sensors mounted at the head of the aircraft
- MCAS activates if **one of the two** AOA sensors says the AOA is too high.

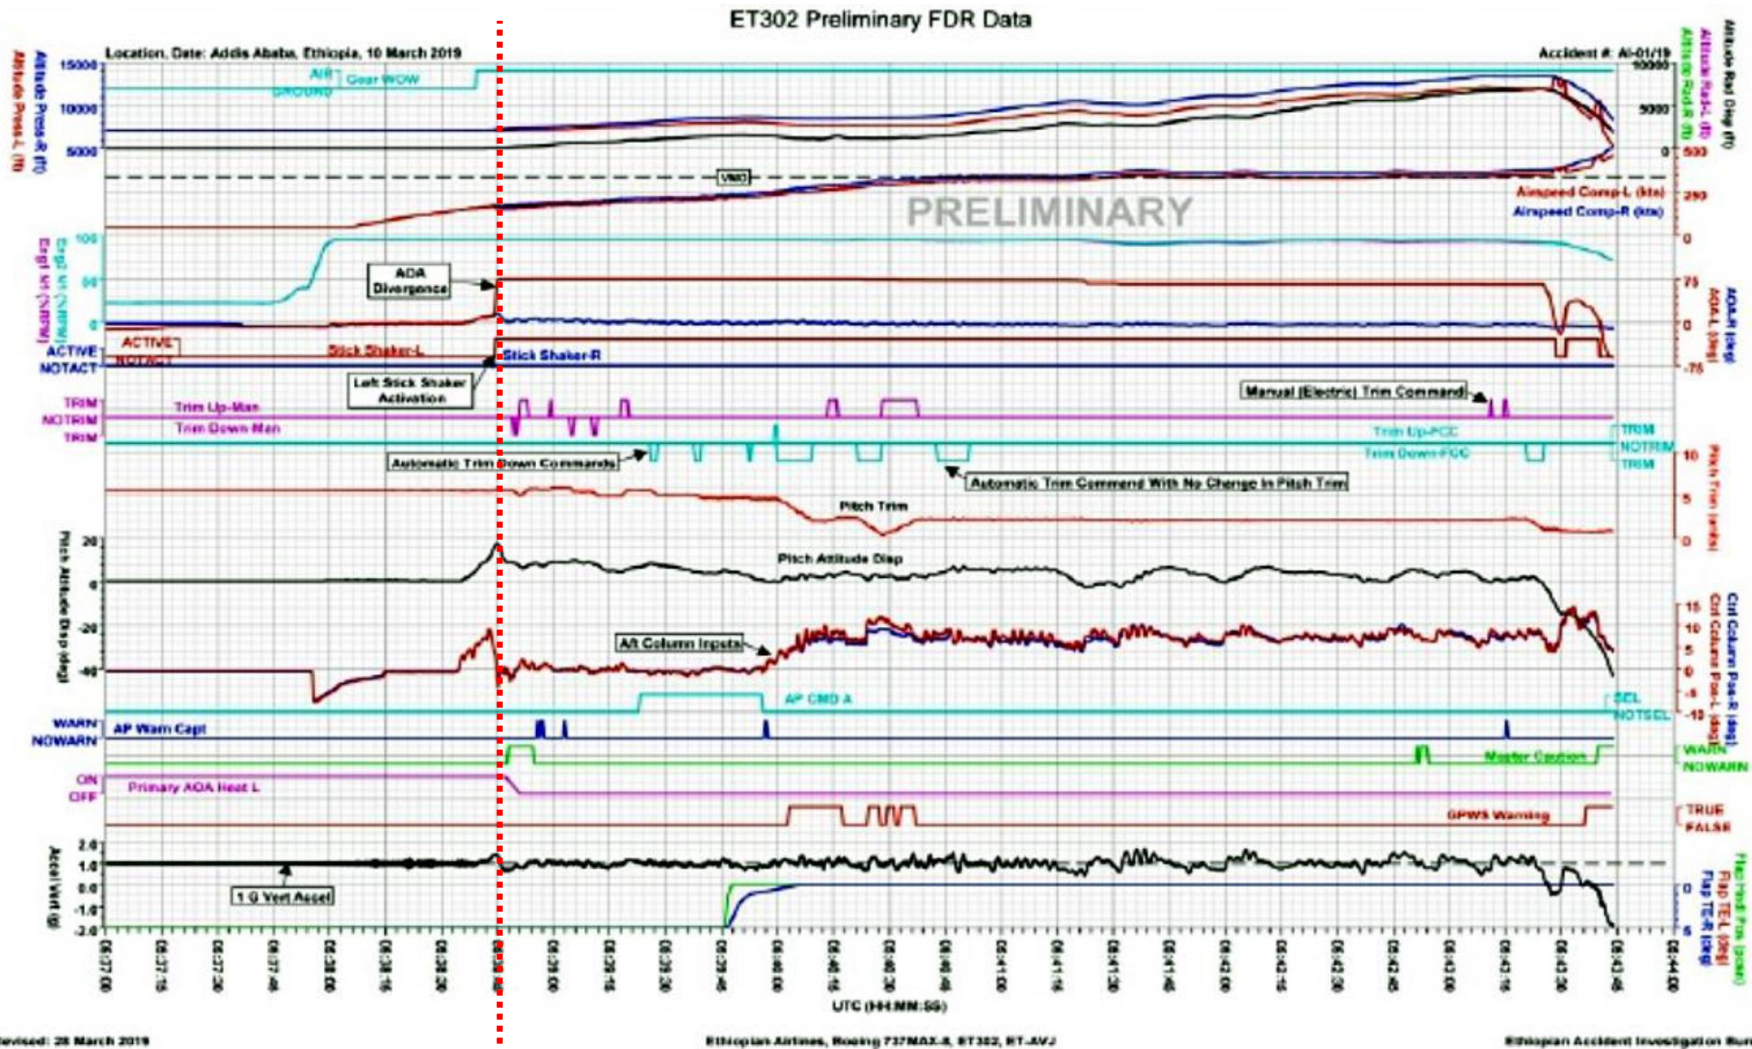


Lion JT610 Flight Data Recording





Sources: Indonesian safety regulators, black box flight recorder data

ET302 Flight Data Recording




How to override MCAS?

- It's not easy...
- Provided instructions to the pilots on how to turn off MCAS



- Available from library website



Lecture 9: Introduction to Model Checking

Formal Verification/Validation



Grand challenge:

Automate the process as much as possible !

Analysis Techniques

- Dynamic Analysis (runtime)
 - Execute the system, possibly multiple times with different inputs
 - Check if every execution meets the desired requirement
- Static Analysis (design time)
 - Analyze the source code or the model for possible bugs
- Trade-offs
 - Dynamic analysis is incomplete, but accurate (checks real system, and bugs discovered are real bugs)
 - Static analysis can catch design bugs early !
 - Many static analysis techniques are not **scalable** (solution: analyze approximate versions, can lead to false warnings)

Verification Methods

- Simulation
 - Simulate the model, possibly multiple times with different inputs
 - Easy to implement, scalable, but no correctness guarantees
- Proof based
 - Construct a proof that system satisfies the invariant
 - Requires manual effort (partial automation possible)
- State-space analysis (Model checking)
 - Algorithm explores “all” reachable states to check invariants
 - Not scalable, but current tools can analyze many real-world designs (relies on many interesting theoretical advances)


Different Requirements

- Safety
 - A system always stays within “good” states (i.e. a nothing bad ever happens)
 - Leader election: it is never the case that two nodes consider them to be leaders
 - Collision avoidance: Distance between two cars is always greater than some minimum threshold
- Liveness
 - System eventually attains its goal
 - Leader election: Each node eventually makes a decision
 - Cruise controller: Actual speed eventually equals desired speed
 - A car will always eventually reach its destination

Model Checking

Model Checking


- A domain D representing the state space of a model
- The reachable state space M for the model
- Define a subset of the state space as property φ
- Explore the whole reachable state space of a model for property violations



Plato vs. Diogenes


- The definition of “*human*”

φ



All living creatures D

Counter – example



Plato

Featherless
Biped

M



Here's Plato's
human!!!!




Diogenes

Challenge

- State space explosion
 - Not every model can be model checked!!
 - i.e. Real-number (continuous) state space
 - Complex dynamics between states
- Solution
 - Simple yet expressive formalisms
 - Symbolic states/executions
 - Model abstraction/approximation


Simple yet expressive formalisms

Basic Finite State Machine (FSM)



WANT: if press is issued twice **quickly** then the **light** will get **brighter**; otherwise the light is turned **off**.

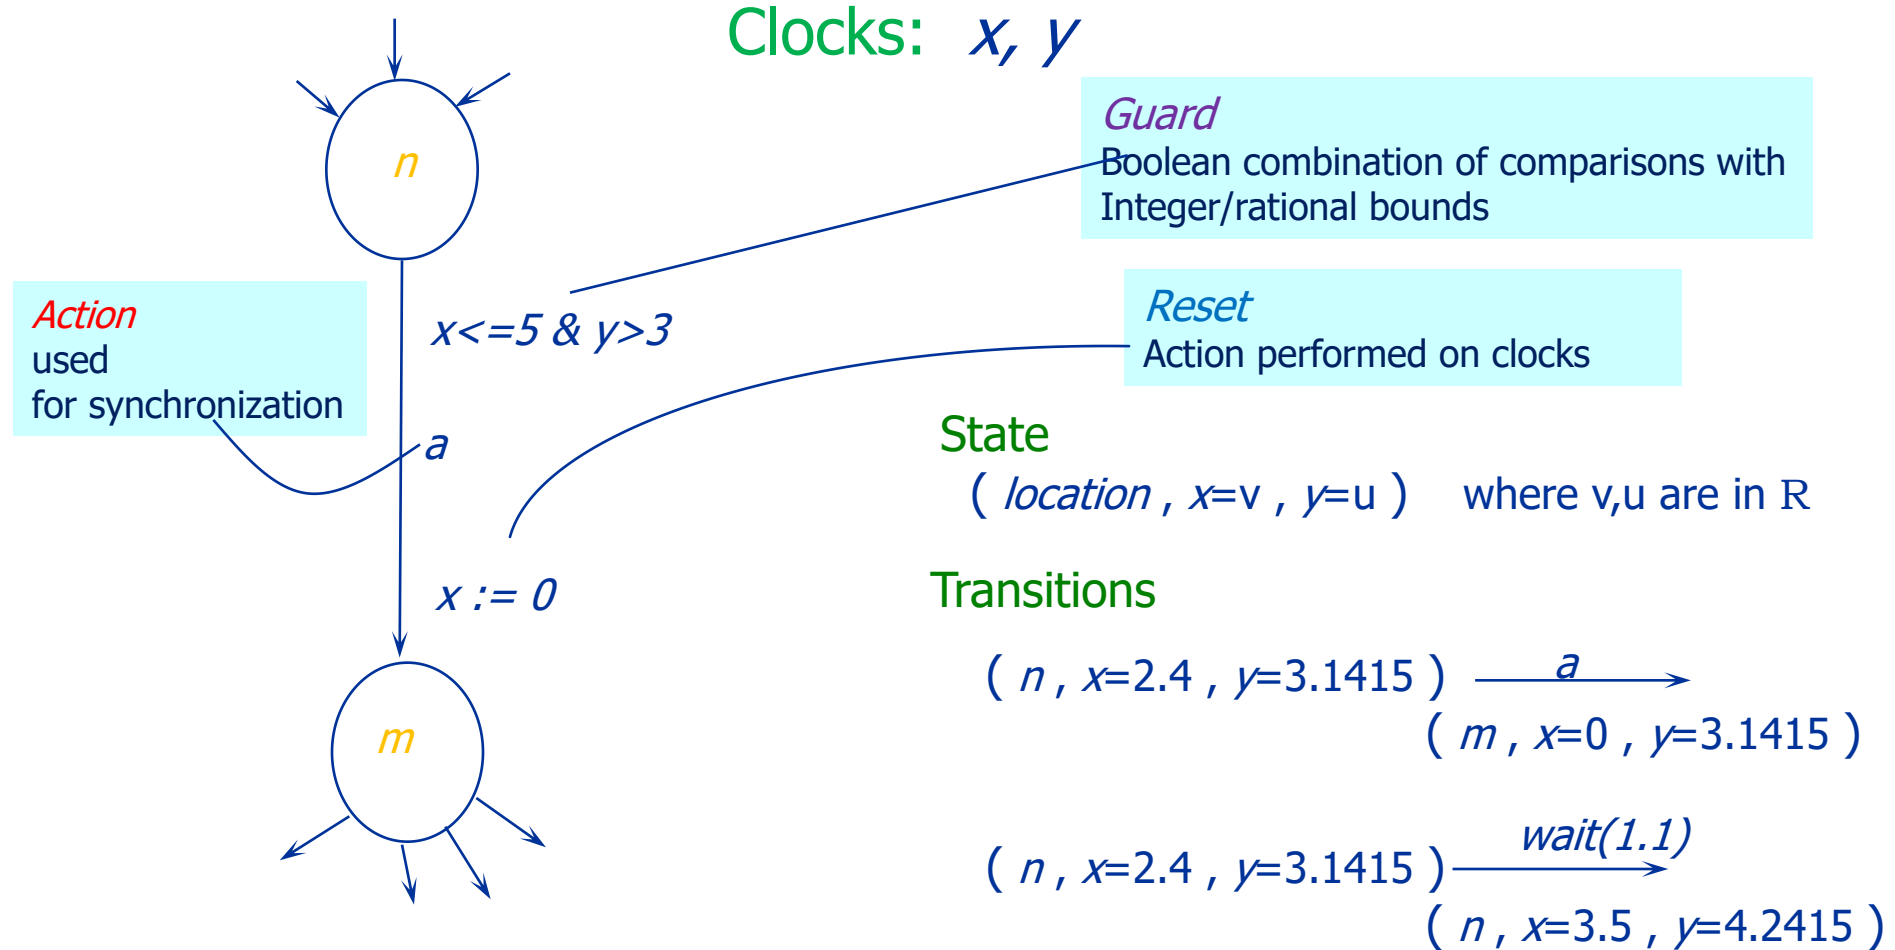
FSM with real number time




Solution: Add a real-valued clock x

Adding continuous variables to state machines

Timed Automata



Adding Invariants



Clocks: x, y


Transitions

$(n, x=2.4, y=3.1415) \xrightarrow{\text{wait}(3.2)}$
 $(n, x=2.4, y=3.1415) \xrightarrow{\text{wait}(1.1)} (n, x=3.5, y=4.2415)$

Invariants ensure progress!!

Timed Automata: Syntax

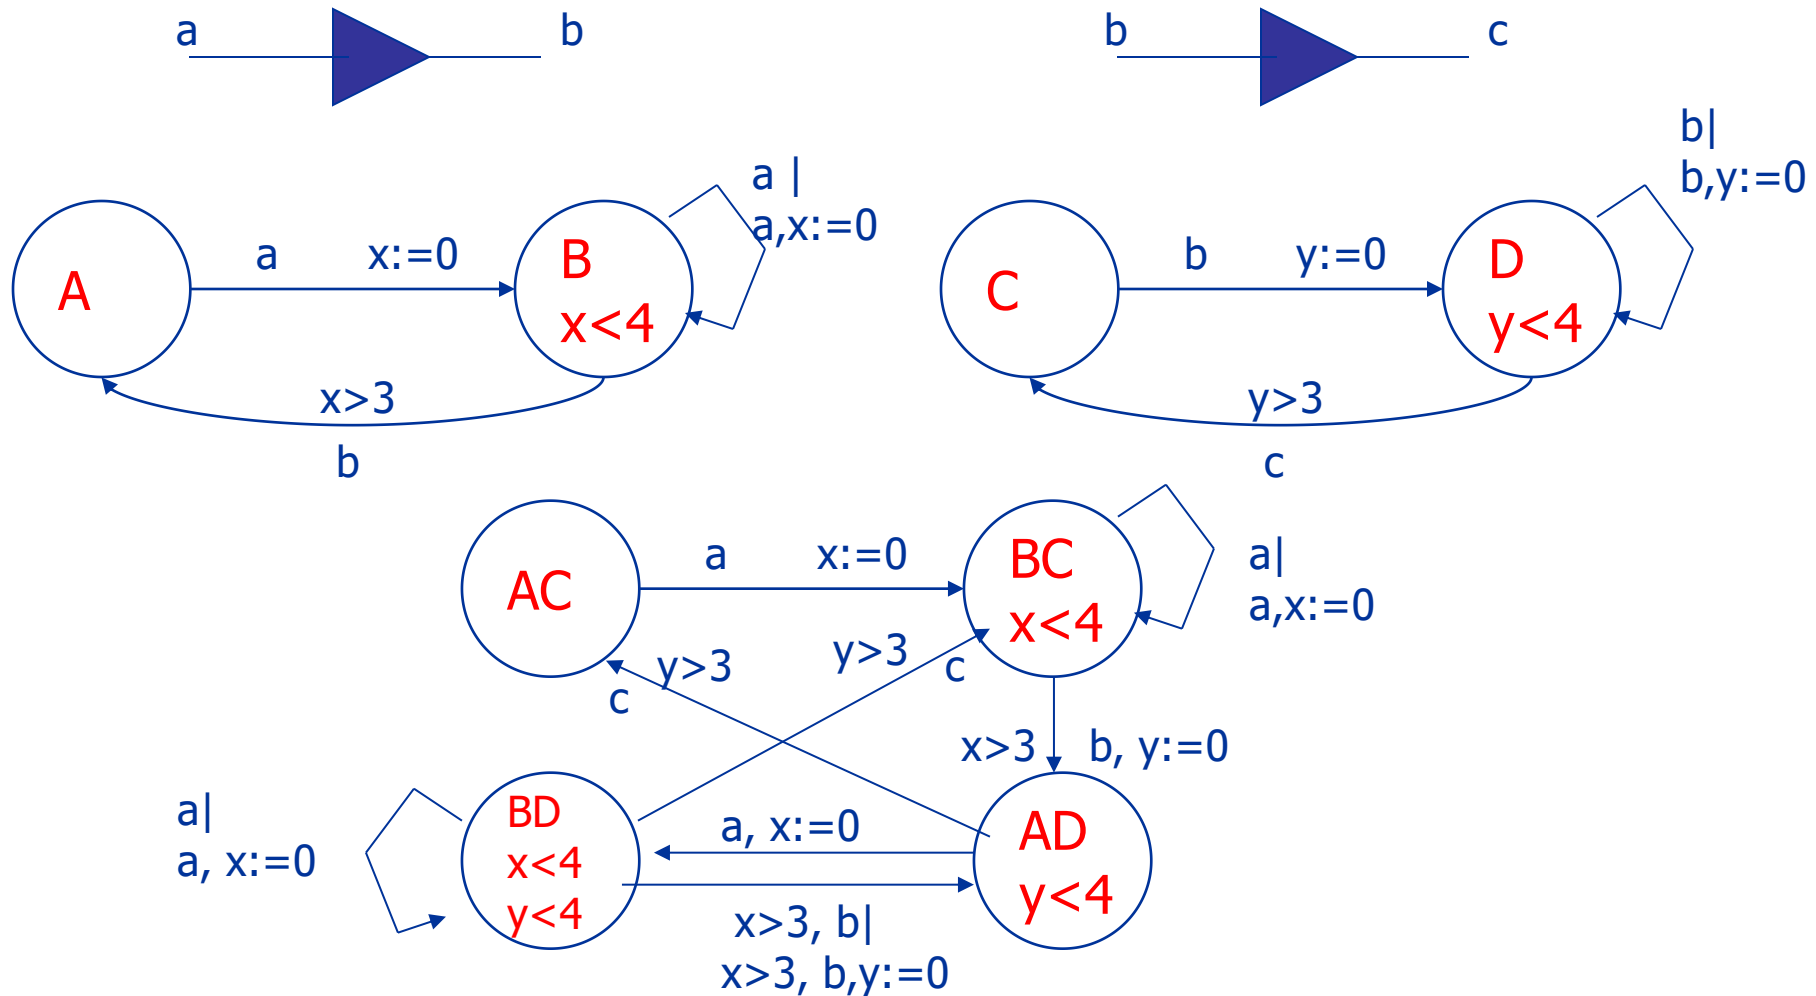
- A finite set V of locations
- A subset V^0 of initial locations
- A finite set Σ of labels (alphabet)
- A finite set X of clocks
- Invariant $Inv(l)$ for each location: (clock constraint over X)
- A finite set E of edges. Each edge has
 - source location l , target location l'
 - label a in Σ (ε labels also allowed)
 - guard g (a clock constraint over X)
 - a subset λ of clocks to be reset



Timed Automata: Semantics

- For a timed automaton A , define an infinite-state transition system $S(A)$
- **States Q :** a state q is a pair (l, v) , where l is a location, and v is a clock vector, mapping clocks in X to R , satisfying $Inv(l)$
- (l, v) is initial state if l is in V^0 and $v(x)=0$
- **Elapse of time transitions:** for each nonnegative real number d , $(l, v) \xrightarrow{d} (l, v+d)$ if both v and $v+d$ satisfy $Inv(l)$
- **Location switch transitions:** $(l, v) \xrightarrow{a} (l', v')$ if there is an edge (l, a, g, λ, l') such that v satisfies g and $v' = v[\lambda := 0]$


Product Construction




Model Checking: Forward Reachability

- Given a timed automata and a property φ
- $R := I$
- Repeat
 - If R intersects $\neg\varphi$, report “yes”
 - Else if R contains $\text{Post}(R)$, report “no”
 - Else $R := R \text{ union } \text{Post}(R)$

Reachability for Timed Automata




gives rise to the
infinite transition system:




Symbolic states/executions

Finite Partitioning

Goal: To partition state-space into finitely many equivalence classes so that equivalent states exhibit similar behaviors




Symbolic States/executions




Model abstraction/approximation

Existential Abstraction (Over-approximation)

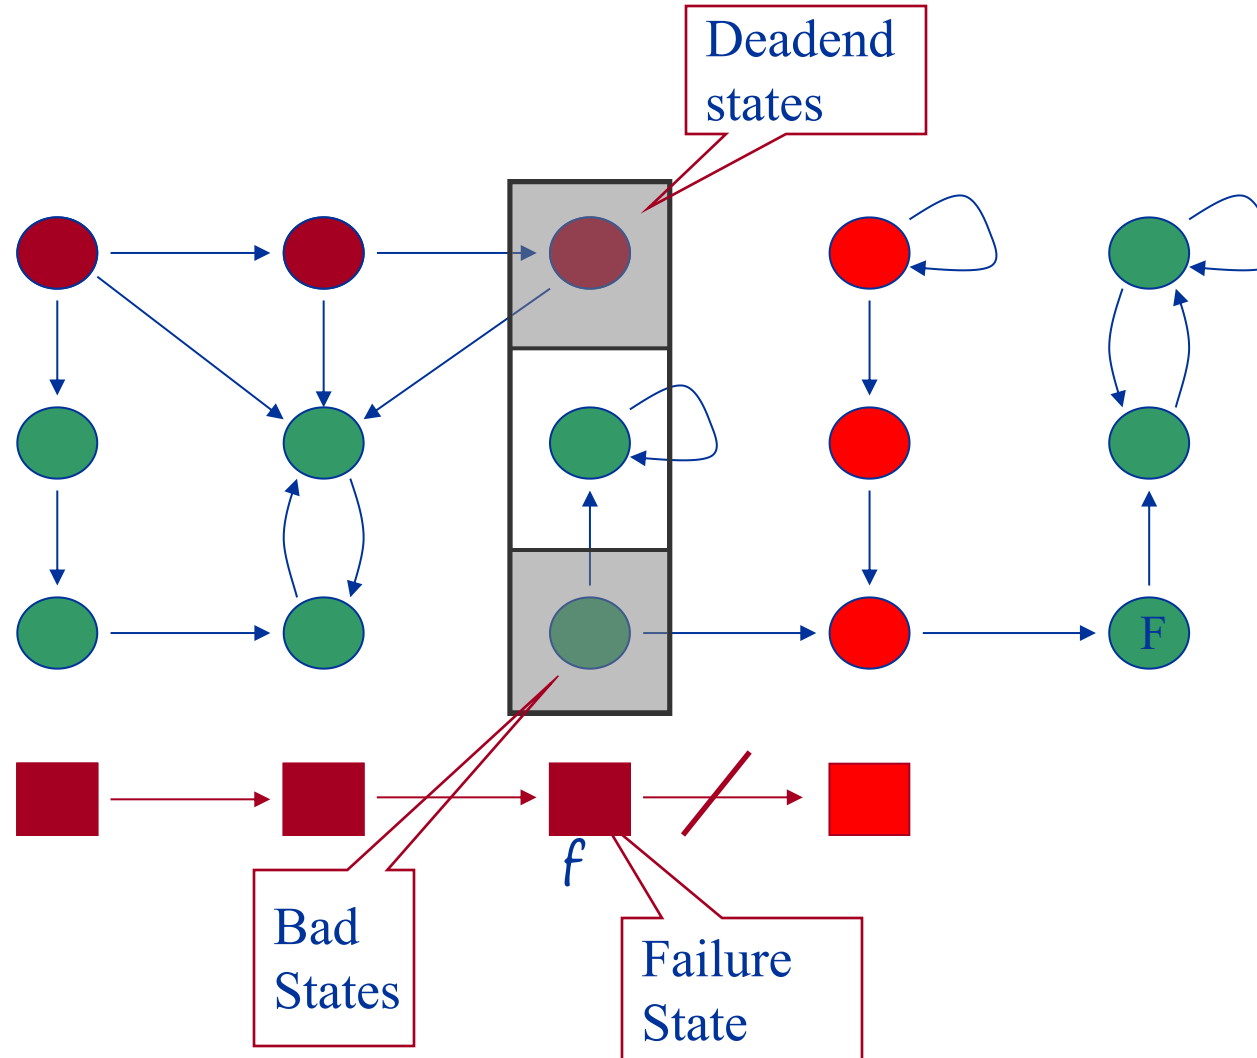


Pros and Cons of Over-approximation

- Properties satisfied by M are also satisfied by P
 - Can model check a less complex model
- M has more behaviors than P
- If a counter-example returns, it may not be a behavior of P




Why spurious counterexample?



Refinement

- **Problem:** Deadend and Bad States are in the same abstract state.
- **Solution:** Refine abstraction function.
- The sets of Deadend and Bad states should be **separated** into different abstract states.

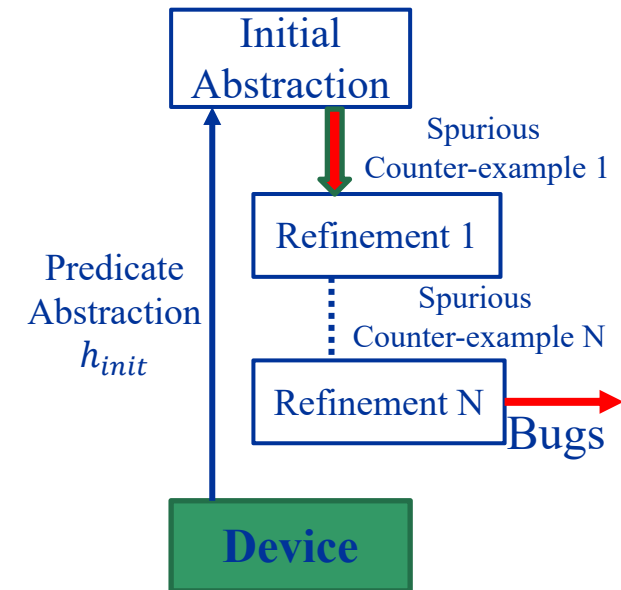
Refinement



Refinement : h'

Counter-Example-Guided Abstraction and Refinement (CEGAR)

- Obtain initial abstraction
 1. Model checking
 2. Property satisfied \rightarrow no bugs
 3. Property unsatisfied \rightarrow counter-examples
 4. Check whether the CE is spurious
 5. If not, bug found
 6. If yes, refine the model and start from 1 again



Capture Environmental Variability With Over-approximation

- Properties satisfied by M are also satisfied by $P1, P2$
- Behaviors not exist in $P1, P2$ may also be physiologically-valid
- Is this a valid counter-example?
- Need a framework to provide context for counter-examples

